

Predicting Heart Disease through Machine Learning

Faadil Ahmed | Audrey Bennett | Shishir Rajendrakumar Deshpande

MSDS-422 Summer 2025

Northwestern University, Practical Machine Learning

July 27, 2025

1. Research Objectives

Our current diagnostic approaches to heart disease are delayed, expensive, and reactive. This not only compromises patient outcomes but also burdens healthcare systems with avoidable costs. Through our analysis, we aim to construct an interpretable machine learning model to predict the likelihood of heart disease and support data-driven early screening and intervention. Further, we aim to integrate our approach with current physician workflows and generate measurable economic value.

2. Exploratory Data Analysis

We analyzed 918 patient observations from the “Kaggle: Heart Disease Prediction” dataset (Janosi et al., 1989). Our target variable, `HeartDisease`, showed reasonable balance with 55% patients with no heart disease and 45% with heart disease. We evaluated the data's distribution patterns and checked numerical variables for outliers and correlations. Assessing the correlation matrix, we observed a negative correlation between heart strain during exercise and the shape of a key EKG reading (ST slope), and a positive correlation between age and cholesterol levels.

3. Data Preparation

To train our machine learning models, we encoded categorical variables and engineered five new features: (1) `Exercise\_Capacity` (maximum heart rate divided by their age) to simulate cardiovascular efficiency, (2) `CVD\_Risk\_Score` (weighted sum of age, resting blood pressure, and cholesterol levels) to indicate a patient's baseline risk, (3) `BP\_Category`, (4) `Chol\_Category`, and (5) `Age\_Group` to understand risk thresholds. Further, we normalized numerical features to ensure scale-sensitive models could assess relationships with the target.

4. Modeling and Results

We used a Python notebook on a CPU-based system to develop our models, leveraging common, open-source libraries, so that they can be deployed and reproduced across all types of clinical environments. We also used traditional machine learning models to maintain explainability and trained them using stratified train-test splits. Beyond training, we tuned models' hyperparameters to optimize their performance.

We analyzed the data via five classification models: (1) Logistic Regression, as a baseline due to its simplicity and interpretability, (2) Random Forest, to capture non-linear interactions, (3) XGBoost, to evaluate dimensional data, (4) Neural Network, to recognize complex patterns like a human brain, and (5) Support Vector Machine, to handle subtle patterns in the data.

Overall, the Random Forest model performed best as it diagnosed heart disease at 89% accuracy and maintained a strong balance between precision and recall (90% F1 score). Further, it successfully classified patients with heart disease, with an ROC-AUC score of 93%. However, we noted that the data had several data quality issues, such as zero values in cholesterol and blood pressure measurements, that can impact model performance when deployed in production.

Through our analysis, we noted that ST slope, the type of chest pain, maximum heart rate, exercise-induced angina, and cholesterol levels were the top identifiers of underlying heart disease. Further, we observed that certain non-obvious interactions influenced predictions, such as high heart strain during exercise combined with flat EKG reading consistently leading to higher disease risks. We confirmed feature importances through Shapley Additive Explanations and revealed that resting blood pressure and age impacted individual predictions the most.

We estimate that deploying this model will save approximately \$3.68 million annually and generate a 400% return on investment through downstream treatment prevention.

5. Technical Considerations

To ensure prediction reliability, we evaluated our models using feature importance tests, cross-validated them on various training data variations, and assessed if our predictions matched our historical and general understanding of heart disease. As a result, our prediction pipeline captures all preprocessing steps, from feature engineering through normalization, and maintains consistency between training and inference.

Additionally, our modeling allows healthcare organizations the flexibility to integrate the model into existing workflows regardless of their technical infrastructure as it supports multiple deployment options, such as REST APIs through Flask/FastAPI or serverless functions via AWS Lambda and Google Cloud Functions.

6. Lessons Learned and Recommendations

Our analysis concentrated on predicting heart disease and understanding the implications of integrating machine learning into clinical settings. Random Forest performed best as the training dataset included under 1,000 observations and a mix of categorical and numerical predictors. In this context, the model's structure allowed it to model complex interactions without overfitting, which is generally a challenge with more data-hungry models like neural networks.

Further, we identified several success factors through this process. First, data quality and governance remain foundational requirements as even small errors could cause harmful clinical recommendations. Second, interpretability matters more than accuracy as physicians cannot stand by results they cannot verify through domain expertise. Third, modeling success will be defined by human factors than technical ones as outcomes will depend on physician decision-making and patient compliance.

Additionally, we recommend that organizations improve data collection practices and improve engagement with physicians to ensure proper training and adoption. We also recommend phased rollouts of all models to minimize churn and allow iterative development through real-world feedback.

In the long term, we recommend that the organizations consider revamping success metrics to go beyond accuracy and into patient outcomes. While technical development is important, we believe that it must be balanced with stakeholder alignment. When clinical trust is paramount, explainability should take precedence over marginal gains in predictive performance.

References

Janosi, Andras, William Steinbrunn, Matthias Pfisterer, and Robert Detrano. 1989. Heart Disease. UCI Machine Learning Repository. <https://doi.org/10.24432/C52P4X>.

Appendix 1: Python Code and Output

Management/Research Question:

What factors most significantly influence the likelihood of heart disease, and how can we predict it using machine learning models?

Our hypothesis is that heart disease can be predicted based on a combination of physiological metrics (e.g., cholesterol levels, age, blood pressure) and categorical health indicators (e.g., type of chest pain, exercise-induced angina, ECG results). To test this, we explore the distribution of the target variable `HeartDisease` and its relationship with both numerical and categorical features. We then build and evaluate multiple classification models including Random Forest, XGBoost, and SVM, with performance comparison and interpretability techniques such as SHAP and permutation importance.

This analysis helps patients and healthcare providers understand key risk factors contributing to heart disease and allows for early intervention. Additionally, it provides a foundation for clinical decision support systems that optimize care and reduce costs.

Methodological Assumptions and Limitations of the Dataset

Assumptions:

- Binary classification (0 = no heart disease, 1 = heart disease) is an appropriate modeling framework.
- All input features are correctly encoded and scaled for model interpretability and performance.
- The dataset is representative of a broader population, and class imbalance is minimal.
- Label encoding adequately captures the influence of categorical health indicators.

Limitations:

- Lifestyle variables (e.g., diet, exercise habits) are not included and could affect prediction accuracy.
- The dataset is cross-sectional and does not capture temporal changes in patient health.
- Some features such as cholesterol and blood pressure may have been measured inconsistently across patients.

EXECUTIVE SUMMARY & RECOMMENDATIONS

MODEL PERFORMANCE SUMMARY:

- **Best Model:** Random Forest
 - **Accuracy:** 0.891
 - **Precision:** 0.894
 - **Recall (Sensitivity):** 0.912
 - **F1-Score:** 0.903
 - **AUC-ROC:** 0.933

BUSINESS IMPACT:

- **Estimated Annual Savings:** \$3,679,500.00
- **Cost per Prediction:** \$5002.72
- **ROI:** 399.7%

These values highlight that deploying the model at scale can result in meaningful financial and clinical impact by preventing downstream costs associated with untreated cardiovascular disease (CVD).

KEY INSIGHTS:

- **Top Risk Factors:** `ST_Slope_encoded`, `ChestPainType_encoded`, `MaxHR`
- **Model Sensitivity:** 0.912 — model is highly sensitive in detecting patients at risk
- **Model Specificity:** 0.866 — also shows strong ability to identify patients not at risk

The combination of high recall and precision makes the model reliable for real-world clinical use where minimizing false negatives is crucial.

RECOMMENDATIONS:

1. Deploy model in **Clinical Decision Support Systems (CDSS)** to assist clinicians.
2. Prioritize interventions for **patients with high CVD risk scores**.
3. Set up **continuous monitoring** for model drift over time.
4. Gather additional **lifestyle data** to enhance prediction accuracy.
5. Integrate seamlessly with **Electronic Health Records (EHRs)**.

NEXT STEPS:

1. **Validate** the model on an external dataset to ensure generalizability.
2. Conduct a **clinical trial** for real-world, prospective validation.
3. Build a **clinician-friendly user interface** for integration into workflows.
4. Define and schedule a **model retraining strategy**.
5. Prepare for **multi-institutional deployment** at scale.

WHY THIS MATTERS:

In healthcare, model transparency, trust, and fairness are critical. This pipeline not only demonstrates strong predictive performance but also offers interpretability through SHAP analysis. This is crucial for:

- Making **transparent, explainable decisions**
- Gaining **clinical trust**
- Identifying hidden correlations between medical features and outcomes
- Supporting **fairer** healthcare deployment

Section 1: SETUP AND IMPORTS

```
In [20]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')

# Machine Learning Libraries
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV, StratifiedKFold
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score, roc_curve
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
import xgboost as xgb

# Model Interpretation
from sklearn.inspection import permutation_importance

# Set display options
pd.set_option('display.max_columns', None)
plt.style.use('seaborn-v0_8')
sns.set_palette("husl")

print("All libraries imported successfully!")
print("Heart Disease Prediction ML Pipeline")
print("Executive Data Science Project")
```

All libraries imported successfully!
Heart Disease Prediction ML Pipeline
Executive Data Science Project

SECTION 2: DATA LOADING AND INITIAL EXPLORATION

```
In [21]: # Load the dataset
df = pd.read_csv('heart.csv')

print("DATASET OVERVIEW")
print("=" * 50)
print(f"Dataset shape: {df.shape}")
```



```
print(f"Memory usage: {df.memory_usage(deep=True).sum() / 1024:.2f} KB")

# Basic info
print("\nCOLUMN INFORMATION")
print("=" * 50)
df.info()

print("\nFIRST 5 ROWS")
print("=" * 50)
display(df.head())

print("\nSTATISTICAL SUMMARY")
print("=" * 50)
display(df.describe(include='all'))
```

DATASET OVERVIEW
=====

Dataset shape: (918, 12)
Memory usage: 281.36 KB

COLUMN INFORMATION
=====

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to 917
Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	Age	918 non-null	int64
1	Sex	918 non-null	object
2	ChestPainType	918 non-null	object
3	RestingBP	918 non-null	int64
4	Cholesterol	918 non-null	int64
5	FastingBS	918 non-null	int64
6	RestingECG	918 non-null	object
7	MaxHR	918 non-null	int64
8	ExerciseAngina	918 non-null	object
9	Oldpeak	918 non-null	float64
10	ST_Slope	918 non-null	object
11	HeartDisease	918 non-null	int64

dtypes: float64(1), int64(6), object(5)
memory usage: 86.2+ KB

FIRST 5 ROWS
=====

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	40	M	ATA	140	289	0	Normal	172	N	0.0	Up	0
1	49	F	NAP	160	180	0	Normal	156	N	1.0	Flat	1
2	37	M	ATA	130	283	0	ST	98	N	0.0	Up	0
3	48	F	ASY	138	214	0	Normal	108	Y	1.5	Flat	1
4	54	M	NAP	150	195	0	Normal	122	N	0.0	Up	0

STATISTICAL SUMMARY
=====

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
count	918.000000	918	918	918.000000	918.000000	918.000000	918	918.000000	918	918.000000	918	918.000000
unique	NaN	2	4	NaN	NaN	NaN	3	NaN	2	NaN	3	NaN
top	NaN	M	ASY	NaN	NaN	NaN	Normal	NaN	N	NaN	Flat	NaN
freq	NaN	725	496	NaN	NaN	NaN	552	NaN	547	NaN	460	NaN
mean	53.510893	NaN	NaN	132.396514	198.799564	0.233115	NaN	136.809368	NaN	0.887364	NaN	0.553377
std	9.432617	NaN	NaN	18.514154	109.384145	0.423046	NaN	25.460334	NaN	1.066570	NaN	0.497414
min	28.000000	NaN	NaN	0.000000	0.000000	0.000000	NaN	60.000000	NaN	-2.600000	NaN	0.000000
25%	47.000000	NaN	NaN	120.000000	173.250000	0.000000	NaN	120.000000	NaN	0.000000	NaN	0.000000
50%	54.000000	NaN	NaN	130.000000	223.000000	0.000000	NaN	138.000000	NaN	0.600000	NaN	1.000000
75%	60.000000	NaN	NaN	140.000000	267.000000	0.000000	NaN	156.000000	NaN	1.500000	NaN	1.000000
max	77.000000	NaN	NaN	200.000000	603.000000	1.000000	NaN	202.000000	NaN	6.200000	NaN	1.000000

SECTION 3: TARGET VARIABLE ANALYSIS

```
In [22]: # Target distribution
target_counts = df['HeartDisease'].value_counts()
target_percent = df['HeartDisease'].value_counts(normalize=True) * 100

print("Heart Disease Distribution:")
for i, (count, percent) in enumerate(zip(target_counts, target_percent)):
    status = "No Heart Disease" if i == 0 else "Heart Disease"
    print(f"    {status}: {count} ({percent:.1f}%)")

# Visualization
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 5))



# Count plot
target_counts.plot(kind='bar', ax=ax1, color=['lightblue', 'lightcoral'])
ax1.set_title('Heart Disease Distribution')
ax1.set_xlabel('Heart Disease (0=No, 1=Yes)')
ax1.set_ylabel('Count')
ax1.tick_params(axis='x', rotation=0)

# Pie chart
ax2.pie(target_counts.values, labels=['No Heart Disease', 'Heart Disease'],
        autopct='%1.1f%%', colors=['lightblue', 'lightcoral'])
ax2.set_title('Heart Disease Proportion')

plt.tight_layout()
plt.show()

# Check for class imbalance
imbalance_ratio = target_counts.min() / target_counts.max()
print(f"\nClass Balance Ratio: {imbalance_ratio:.3f}")
if imbalance_ratio < 0.5:
    print("    Dataset is imbalanced - consider using stratified sampling")
else:
    print("    Dataset is reasonably balanced")
```

Heart Disease Distribution:
No Heart Disease: 508 (55.3%)
Heart Disease: 410 (44.7%)



Class Balance Ratio: 0.807
Dataset is reasonably balanced

TARGET VARIABLE ANALYSIS

This section investigates the distribution of the **target variable**, `HeartDisease`, which is binary:

- 0 = No heart disease

- 1 = Heart disease present

1. Distribution Summary

We calculate both the **raw counts** and **percentage** of each class:

- 508 patients (55.3%) do **not** have heart disease
- 410 patients (44.7%) **do** have heart disease

This near-even split suggests we won't face major class imbalance issues in model training.

2. Visual Analysis

Two plots provide a quick visual summary:

- **Bar Plot:** Emphasizes frequency of both classes.
- **Pie Chart:** Shows proportional representation (slightly more class 0).

These visuals confirm the numeric results and help assess the need for resampling techniques.

3. Class Imbalance Check

We compute the **class balance ratio**:

smallest class size / largest class size = 0.807

A ratio above 0.5 indicates a **reasonably balanced** dataset. This means:

- No need for SMOTE or undersampling
- Standard evaluation metrics like accuracy and F1-score remain reliable

Conclusion:

No imbalance mitigation techniques are needed at this stage. The class distribution is suitable for supervised classification without additional correction.

SECTION 4: COMPREHENSIVE EDA

```
In [23]: # Numerical features analysis
numerical_cols = df.select_dtypes(include=[np.number]).columns.tolist()
categorical_cols = df.select_dtypes(include=['object']).columns.tolist()

print(f"Numerical features: {numerical_cols}")
print(f"Categorical features: {categorical_cols}")
print(f"Number of categorical features: {len(categorical_cols)}")


# 4.1 Numerical Features Distribution
fig, axes = plt.subplots(2, 3, figsize=(18, 12))
axes = axes.ravel()

for i, col in enumerate(numerical_cols[:-1]): # Exclude target variable
    df[col].hist(bins=30, ax=axes[i], alpha=0.7, color='skyblue', edgecolor='black')
    axes[i].set_title(f'Distribution of {col}')
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('Frequency')

    # Add mean line
    axes[i].axvline(df[col].mean(), color='red', linestyle='--',
                    label=f'Mean: {df[col].mean():.1f}')
    axes[i].legend()

plt.tight_layout()
plt.show()

Numerical features: ['Age', 'RestingBP', 'Cholesterol', 'FastingBS', 'MaxHR', 'Oldpeak', 'HeartDisease']
Categorical features: ['Sex', 'ChestPainType', 'RestingECG', 'ExerciseAngina', 'ST_Slope']
Number of categorical features: 5
```



In this step, we begin exploratory data analysis by examining the distribution of **numerical features** in the dataset. This helps us understand data ranges, detect skewness, identify anomalies or outliers, and guide potential transformations.

1. Feature Type Breakdown

The dataset includes:

- **7 numerical features:** Age, RestingBP, Cholesterol, FastingBS, MaxHR, Oldpeak, HeartDisease
- **5 categorical features:** Sex, ChestPainType, RestingECG, ExerciseAngina, ST_Slope

The target variable `HeartDisease` is also numerical (binary), but we exclude it from the EDA distribution analysis at this stage.

2. Histogram Visualization

Each numerical feature is plotted as a histogram with a red vertical line showing the **mean**. Observations:

- **Age**: Appears roughly normally distributed, centered around age ~53.
- **RestingBP**: Fairly concentrated around 120–140, but a few entries fall below physiological norms (possible data issues).
- **Cholesterol**: Right-skewed; some extremely high values (>500) may indicate outliers or entry errors.
- **FastingBS**: Highly imbalanced; majority have fasting blood sugar ≤ 120 mg/dL (**0**), few have high (**1**). Will be treated as binary.
- **MaxHR**: Slight left-skew, with a mean around 136 bpm. No major outliers.
- **Oldpeak**: Strong right skew. Most values are 0, indicating minimal ST depression during exercise — only a few patients have large values.

3. Why This Matters

Understanding feature distributions:

- Informs **normalization or transformation** decisions (e.g., log-transform skewed variables)
- Helps with **outlier detection** and handling
- Highlights possible **data quality issues** (e.g., zero cholesterol or blood pressure)
- Aids in **feature selection** for modeling

In [24]:


```
# 4.2 Categorical Features Analysis
# Calculate number of rows and columns needed
num_categorical = len(categorical_cols)
cols = 3 # 3 columns
rows = (num_categorical + cols - 1) // cols # Calculate rows needed

fig, axes = plt.subplots(rows, cols, figsize=(15, 5 * rows))
if rows == 1:
    axes = [axes] # Handle single row case
axes = axes.ravel()

for i, col in enumerate(categorical_cols):
    value_counts = df[col].value_counts()
    axes[i].bar(value_counts.index, value_counts.values, color='lightgreen', alpha=0.7)
    axes[i].set_title(f'Distribution of {col}')
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('Count')
    axes[i].tick_params(axis='x', rotation=45)

# Hide any unused subplots
for i in range(num_categorical, len(axes)):
    axes[i].set_visible(False)

plt.tight_layout()
plt.show()
```



After examining the numerical features, we now explore the **categorical variables** in the dataset to understand the frequency distribution of each category. This helps identify class dominance, imbalance, or data encoding issues.

1. Setup and Visualization

There are **5 categorical columns** in the dataset:

- Sex
- ChestPainType
- RestingECG
- ExerciseAngina
- ST_Slope

A bar plot is created for each variable using `value_counts()`. These charts show how often each category appears in the dataset.

2. Observations by Feature

- **Sex:** Skewed toward males, with ~725 male and ~193 female entries. Gender may need encoding and possibly balancing if found predictive.
- **ChestPainType:** Dominated by `ASY` (asymptomatic) cases, followed by `NAP`, `ATA`, and a small number of `TA`. This imbalance could impact model learning and feature importance.
- **RestingECG:** Most patients have a `Normal` reading, with fewer in `LVH` (left ventricular hypertrophy) or `ST`.
- **ExerciseAngina:** More patients answered `No` (547) than `Yes` (371), which could influence risk scoring.
- **ST_Slope:** `Flat` and `Up` are the most common ST slope patterns, while `Down` is rare — possibly a strong signal if predictive.

💡 3. Why This Matters

- These distributions reveal **dominant classes** and **potential features for one-hot encoding**.
- Rare categories (like `TA` in `ChestPainType` or `Down` in `ST_Slope`) may be strong predictors if associated with high disease incidence.
- If category imbalance is too severe, encoding strategies (e.g., target encoding) or grouping rare values may be needed.

Conclusion:

Understanding how categorical features are distributed sets the foundation for encoding, interaction analysis, and feature engineering in upcoming steps.

In [25]:


```
# 4.3 Correlation Analysis
print("\nCORRELATION ANALYSIS")
print("=" * 30)

# Correlation matrix for numerical features
correlation_matrix = df[numerical_cols].corr()

plt.figure(figsize=(10, 8))
mask = np.triu(np.ones_like(correlation_matrix, dtype=bool))
sns.heatmap(correlation_matrix, mask=mask, annot=True, cmap='coolwarm', center=0,
            square=True, fmt='.2f', cbar_kws={"shrink": .8})
plt.title('Correlation Matrix - Numerical Features')
plt.tight_layout()
plt.show()

# Strong correlations with target
target_corr = correlation_matrix['HeartDisease'].abs().sort_values(ascending=False)
print("\nCorrelations with Heart Disease (strongest first):")
for feature, corr in target_corr.items():
    if feature != 'HeartDisease':
        print(f"  {feature}: {corr:.3f}")
```

CORRELATION ANALYSIS
=====



Correlations with Heart Disease (strongest first):

Oldpeak: 0.404
MaxHR: 0.400
Age: 0.282
FastingBS: 0.267
Cholesterol: 0.233
RestingBP: 0.108

This section evaluates the relationships between numerical features and the target variable `HeartDisease` using a correlation matrix.

- **correlation_matrix**: Computes Pearson correlation coefficients between numerical features.
- **mask = np.triu(...)**: Hides the upper triangle of the heatmap to reduce redundancy.
- **sns.heatmap(...)**: Visualizes the correlation matrix using a diverging color map (`coolwarm`), centered at 0 to highlight positive vs negative correlations.
- **annot=True**: Displays the numeric correlation values within each cell.
- **cbar_kws={"shrink": .8}**: Adjusts the size of the color bar for better layout.

Below the heatmap, the code extracts and displays the absolute correlation values of each feature with the target `HeartDisease` , sorted in descending order (strongest relationships first):

- **target_corr = correlation_matrix['HeartDisease'].abs().sort_values(...)**: Ranks features by strength of correlation with the target.
- **Excludes self-correlation** (`feature != 'HeartDisease'`).

This analysis helps identify which numerical features are most associated with the presence of heart disease, guiding further feature selection or engineering.

SECTION 5: ADVANCED FEATURE ANALYSIS

```
In [26]: print("\n ADVANCED FEATURE ANALYSIS")
print("=" * 50)

# 5.1 Age analysis
```



```
fig, axes = plt.subplots(1, 3, figsize=(18, 5))

# Age distribution by heart disease
df.boxplot(column='Age', by='HeartDisease', ax=axes[0])
axes[0].set_title('Age Distribution by Heart Disease')
axes[0].set_xlabel('Heart Disease (0=No, 1=Yes)')


# Age histogram by heart disease
for hd in [0, 1]:
    subset = df[df['HeartDisease'] == hd]
    axes[1].hist(subset['Age'], alpha=0.7, label=f'HD={hd}', bins=20)
axes[1].set_title('Age Distribution by Heart Disease Status')
axes[1].set_xlabel('Age')
axes[1].legend()

# Age groups analysis
df['AgeGroup'] = pd.cut(df['Age'], bins=[0, 40, 50, 60, 100],
                        labels=['<40', '40-50', '50-60', '60+'])
age_hd_crosstab = pd.crosstab(df['AgeGroup'], df['HeartDisease'], normalize='index')
age_hd_crosstab.plot(kind='bar', ax=axes[2], stacked=False)
axes[2].set_title('Heart Disease Rate by Age Group')
axes[2].set_ylabel('Proportion')
axes[2].tick_params(axis='x', rotation=45)

plt.tight_layout()
plt.show()
```

ADVANCED FEATURE ANALYSIS

=====



In this section, we take a deeper look at how **age** interacts with heart disease status using multiple visual techniques. These insights can guide feature transformation, binning, and selection.

1. Boxplot: Age Distribution by Heart Disease

```
df.boxplot(column='Age', by='HeartDisease', ax=axes[0])
```

- This shows **age distribution split by heart disease presence** using a boxplot.
- Patients with heart disease (1) generally skew **older** than those without (0).
- This visualization highlights **outliers** and **median age differences**, helping assess if age is a good predictor.

2. Histogram: Overlaid Age Distribution by Heart Disease Status

```
for hd in [0, 1]:
    subset = df[df['HeartDisease'] == hd]
    axes[1].hist(subset['Age'], alpha=0.7, ...)
```

- Overlays two histograms: one for each class (HeartDisease = 0 and 1).
- Clearly shows **age concentration** differences: heart disease is more common in the **50-65** range.
- Helps detect **overlapping regions** where the two classes may be harder to separate.

3. Age Group Analysis: Heart Disease Proportion by Bin

```
df['AgeGroup'] = pd.cut(df['Age'], bins=..., labels=...)
pd.crosstab(..., normalize='index')
```

- Age is binned into 4 ranges: <40 , 40-50 , 50-60 , and 60+ .
- For each group, the **proportion of patients with heart disease** is plotted.
- Heart disease **increases with age**, peaking around 60+ , where over **70%** have the condition.
- `normalize='index'` ensures the chart reflects **relative proportions**, not raw counts.

Key Insights:

- **Age** is clearly **positively correlated** with heart disease and should be treated as a **critical predictor**.
- Consider engineering an `AgeGroup` feature for tree-based models or stratified analysis.
- Visualization confirms that both **distribution** and **binned rate** offer valuable modeling signals.

Conclusion:

These advanced visualizations support **feature importance reasoning**, justify **age-based segmentation**, and guide further preprocessing or model tuning steps.

```
In [27]: # 5.2 Gender analysis
print("\n GENDER ANALYSIS")
gender_analysis = pd.crosstab(df['Sex'], df['HeartDisease'], normalize='index')
print("Heart Disease Rate by Gender:")
print(gender_analysis)

# 5.3 Chest Pain Type analysis
print("\n CHEST PAIN TYPE ANALYSIS")
chest_pain_analysis = pd.crosstab(df['ChestPainType'], df['HeartDisease'], normalize='index')
print("Heart Disease Rate by Chest Pain Type:")
print(chest_pain_analysis)

# Visualization
fig, axes = plt.subplots(1, 2, figsize=(15, 5))


# Gender
gender_analysis.plot(kind='bar', ax=axes[0])
axes[0].set_title('Heart Disease Rate by Gender')
axes[0].set_ylabel('Proportion')
axes[0].tick_params(axis='x', rotation=0)

# Chest Pain Type
chest_pain_analysis.plot(kind='bar', ax=axes[1])
axes[1].set_title('Heart Disease Rate by Chest Pain Type')
axes[1].set_ylabel('Proportion')
axes[1].tick_params(axis='x', rotation=45)

plt.tight_layout()
plt.show()
```

GENDER ANALYSIS		
Heart Disease Rate by Gender:		
HeartDisease	0	1
Sex		
F	0.740933	0.259067
M	0.368276	0.631724

CHEST PAIN TYPE ANALYSIS		
Heart Disease Rate by Chest Pain Type:		
HeartDisease	0	1
ChestPainType		
ASY	0.209677	0.790323
ATA	0.861272	0.138728
NAP	0.645320	0.354680
TA	0.565217	0.434783



SECTION 5.2 & 5.3: GENDER AND CHEST PAIN TYPE ANALYSIS – ANNOTATIONS

In this section, we examine how **gender** and **chest pain type** relate to heart disease using crosstabulations and visualizations.

Gender-Based Heart Disease Analysis

- A normalized crosstab (`pd.crosstab`) calculates the **proportion of heart disease presence (1)** and **absence (0)** within each gender.
- The resulting table shows that:
 - **Females (F)** have a higher proportion of no heart disease (74%) compared to presence (26%).
 - **Males (M)** show the opposite trend — a higher incidence of heart disease (63%) than not (37%).

Why this matters:

The gender disparity suggests males in the dataset may be at a higher risk of heart disease. This could be a useful predictive feature for modeling.

Chest Pain Type and Heart Disease

- Another normalized crosstab analyzes `ChestPainType` against `HeartDisease`.
- Notable observations:
 - **ASY (Asymptomatic)** pain type is heavily skewed toward heart disease (79% positive).
 - **ATA (Atypical Angina)** is strongly associated with **no heart disease** (86% negative).
 - **NAP** and **TA** show a more balanced but informative distribution.

Why this matters:

Chest pain type is one of the strongest categorical indicators of heart disease in this dataset. The stark contrast between `ASY` and `ATA` cases could be critical for early diagnosis models.

Visualizations

Two bar plots are generated:

1. **Heart Disease Rate by Gender:**
 - Visualizes the proportion of heart disease per gender using side-by-side bars.
2. **Heart Disease Rate by Chest Pain Type:**
 - Shows how different chest pain categories correlate with disease outcome.

These visuals help quickly identify categorical patterns and strengthen feature understanding for modeling.

Conclusion:

Both gender and chest pain type show **distinct predictive value** for heart disease. These features should be prioritized during feature engineering and model selection.

SECTION 6: FEATURE ENGINEERING

```
In [28]: # Create a copy for feature engineering
df_engineered = df.copy()

# 6.1 Create new features
print("Creating new features...")

# Cardiovascular Risk Score (simplified)
df_engineered['CVD_Risk_Score'] = (
    (df_engineered['Age'] > 55).astype(int) * 2 +
    (df_engineered['Sex'] == 'M').astype(int) * 1 +
    (df_engineered['Cholesterol'] > 240).astype(int) * 2 +
    (df_engineered['RestingBP'] > 140).astype(int) * 2 +
    (df_engineered['FastingBS'] == 1).astype(int) * 1
)

# Exercise Capacity Index
df_engineered['Exercise_Capacity'] = df_engineered['MaxHR'] / df_engineered['Age']

# BP Category
df_engineered['BP_Category'] = pd.cut(df_engineered['RestingBP'],
                                     bins=[0, 120, 140, 180, 300],
                                     labels=['Normal', 'Elevated', 'High', 'Very High'])

# Cholesterol Category
df_engineered['Chol_Category'] = pd.cut(df_engineered['Cholesterol'],
                                     bins=[0, 200, 240, 1000],
                                     labels=['Normal', 'Borderline', 'High'])

# Age groups
df_engineered['Age_Group'] = pd.cut(df_engineered['Age'],
                                   bins=[0, 45, 60, 100],
                                   labels=['Young', 'Middle', 'Senior'])

print(" New features created:")
new_features = ['CVD_Risk_Score', 'Exercise_Capacity', 'BP_Category', 'Chol_Category', 'Age_Group']
for feature in new_features:
    print(f" - {feature}")

# 6.2 Encode categorical variables
print("\n ENCODING CATEGORICAL VARIABLES")

# Create encoders
label_encoders = {}
categorical_features = ['Sex', 'ChestPainType', 'RestingECG', 'ExerciseAngina', 'ST_Slope',
                       'BP_Category', 'Chol_Category', 'Age_Group']

df_encoded = df_engineered.copy()

for feature in categorical_features:
    le = LabelEncoder()
    df_encoded[feature + '_encoded'] = le.fit_transform(df_encoded[feature])
    label_encoders[feature] = le
    print(f" {feature} encoded")

# 6.3 Feature scaling preparation
print("\n PREPARING FOR FEATURE SCALING")

# Identify features for scaling
numerical_features = ['Age', 'RestingBP', 'Cholesterol', 'MaxHR', 'Oldpeak',
                     'CVD_Risk_Score', 'Exercise_Capacity']

print(f"Features to be scaled: {numerical_features}")
```

Creating new features...

New features created:

- CVD_Risk_Score
- Exercise_Capacity
- BP_Category
- Chol_Category
- Age_Group

ENCODING CATEGORICAL VARIABLES

Sex encoded
ChestPainType encoded
RestingECG encoded
ExerciseAngina encoded
ST_Slope encoded
BP_Category encoded
Chol_Category encoded
Age_Group encoded

PREPARING FOR FEATURE SCALING

Features to be scaled: ['Age', 'RestingBP', 'Cholesterol', 'MaxHR', 'Oldpeak', 'CVD_Risk_Score', 'Exercise_Capacity']

Feature engineering is a critical step to derive more informative variables and transform raw features into model-friendly formats. This section performs three main tasks: new feature creation, encoding categorical variables, and preparing for scaling.

6.1 Creating New Features

Several new features are created based on medical logic and domain knowledge:

- **CVD_Risk_Score:** A simplified cardiovascular risk score is computed using binary indicators for age > 55, sex = male, cholesterol > 240, resting BP > 140, and fasting blood sugar = 1. Each factor contributes weighted points toward the score.
- **Exercise_Capacity:** Calculated as `MaxHR / Age`, this feature estimates the patient's cardiovascular efficiency.
- **BP_Category:** `RestingBP` is binned into categories like "Normal", "Elevated", "High", and "Very High".
- **Chol_Category:** `Cholesterol` is categorized as "Normal", "Borderline", or "High".
- **Age_Group:** Age is bucketed into groups: "Young" (<45), "Middle" (45–60), and "Senior" (>60).

These transformations help the model better interpret non-linear medical risk factors.

6.2 Encoding Categorical Variables

To prepare for machine learning, all categorical variables must be converted to numerical values:

- Categorical features like `Sex`, `ChestPainType`, `RestingECG`, `ExerciseAngina`, `ST_Slope`, `BP_Category`, `Chol_Category`, and `Age_Group` are **label-encoded**.
- Label encoding assigns an integer to each unique category in a feature, enabling models to process them numerically.

A dictionary of encoders is saved for reference or inverse transformation later.

6.3 Preparing for Feature Scaling

Finally, the notebook identifies numerical features to be scaled in the next step:

- Selected features include both raw values (like `Age`, `RestingBP`, `Cholesterol`) and engineered ones (like `CVD_Risk_Score` and `Exercise_Capacity`).
- These features will be normalized using standard scaling techniques to ensure consistent scale and distribution.

Conclusion:

This stage builds a solid foundation for the model by enriching the dataset with clinically relevant insights and converting all features into a usable format. The thoughtful categorization and encoding will likely enhance model accuracy and interpretability.

SECTION 7: MODEL PREPARATION

```
In [29]: # Prepare features and target
feature_columns = (numerical_features +
                  [col for col in df_encoded.columns if col.endswith('_encoded')] +
                  ['FastingBS']) # FastingBS is already binary

X = df_encoded[feature_columns]
y = df_encoded['HeartDisease']

print(f"Feature matrix shape: {X.shape}")
print(f"Target vector shape: {y.shape}")
```

```
print(f"Features used: {list(X.columns)}")

# Split the data
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

print(f"\n Data Split:")
print(f" Training set: {X_train.shape[0]} samples")
print(f" Test set: {X_test.shape[0]} samples")
print(f" Training target distribution: {y_train.value_counts().tolist()}")
print(f" Test target distribution: {y_test.value_counts().tolist()}")

# Scale numerical features
scaler = StandardScaler()
X_train_scaled = X_train.copy()
X_test_scaled = X_test.copy()

X_train_scaled[numerical_features] = scaler.fit_transform(X_train[numerical_features])
X_test_scaled[numerical_features] = scaler.transform(X_test[numerical_features])

print(" Features scaled successfully")
```

Feature matrix shape: (918, 16)

Target vector shape: (918,)

Features used: ['Age', 'RestingBP', 'Cholesterol', 'MaxHR', 'Oldpeak', 'CVD_Risk_Score', 'Exercise_Capacity', 'Sex_encoded', 'ChestPainType_encoded', 'RestingECG_encoded', 'ExerciseAngina_encoded', 'ST_Slope_encoded', 'BP_Category_encoded', 'Chol_Category_encoded', 'Age_Group_encoded', 'FastingBS']

Data Split:

Training set: 734 samples

Test set: 184 samples

Training target distribution: [406, 328]

Test target distribution: [102, 82]

Features scaled successfully

SECTION 7: MODEL PREPARATION – ANNOTATIONS

Before training the machine learning model, we must prepare our dataset by selecting features, splitting the data, and applying standard scaling to the numerical variables.

1. Feature and Target Selection

- `feature_columns` is built by combining:
 - The list of raw numerical features (e.g., Age, Cholesterol, MaxHR, etc.)
 - All label-encoded categorical features (columns ending with `_encoded`)
 - `FastingBS`, which is already binary
- `X` contains the feature matrix; `y` contains the target variable `HeartDisease`.

2. Train-Test Split

- The dataset is split using an 80/20 ratio via `train_test_split`.
- Stratification ensures the proportion of heart disease classes is preserved in both train and test sets.
- Output confirms the number of samples and class distribution in each subset:
 - Training set: 734 samples
 - Test set: 184 samples
 - Balanced target distributions maintained using `stratify=y`

3. Scaling Numerical Features

- Standardization (z-score scaling) is applied to numerical features using `StandardScaler`:
 - `fit_transform()` is applied on the training data.
 - `transform()` is applied to the test data using the same scaler.

This step ensures all numerical features have mean 0 and standard deviation 1, preventing bias toward larger-value features during model training.

Conclusion:

With features selected, data split, and values standardized, the dataset is fully prepared for training machine learning models. This setup ensures consistency, fairness, and robustness during learning and evaluation.

SECTION 8: MODEL DEVELOPMENT AND EVALUATION

```

In [30]: # Initialize models
models = {
    'Logistic Regression': LogisticRegression(random_state=42, max_iter=1000),
    'Random Forest': RandomForestClassifier(random_state=42, n_estimators=100),
    'XGBoost': xgb.XGBClassifier(random_state=42, eval_metric='logloss'),
    'Neural Network': MLPClassifier(random_state=42, max_iter=1000),
    'SVM': SVC(random_state=42, probability=True)
}

# Cross-validation results storage
cv_results = {}
model_performances = {}

print(" CROSS-VALIDATION EVALUATION")
print("-" * 40)

# Perform cross-validation for each model
cv_folds = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

for name, model in models.items():
    print(f"\n Evaluating {name}...")

    # Cross-validation scores
    cv_scores = cross_val_score(model, X_train_scaled, y_train,
                                cv=cv_folds, scoring='accuracy')
    cv_results[name] = cv_scores

    # Fit model and make predictions
    model.fit(X_train_scaled, y_train)
    y_pred = model.predict(X_test_scaled)
    y_pred_proba = model.predict_proba(X_test_scaled)[: , 1] if hasattr(model, 'predict_proba') else None

    # Calculate metrics
    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    auc = roc_auc_score(y_test, y_pred_proba) if y_pred_proba is not None else None

    model_performances[name] = {
        'cv_mean': cv_scores.mean(),
        'cv_std': cv_scores.std(),
        'accuracy': accuracy,
        'precision': precision,
        'recall': recall,
        'f1': f1,
        'auc': auc,
        'model': model
    }

    print(f" CV Accuracy: {cv_scores.mean():.4f} (+/- {cv_scores.std() * 2:.4f})")
    print(f" Test Accuracy: {accuracy:.4f}")
    print(f" Test F1-Score: {f1:.4f}")
    if auc:
        print(f" Test AUC: {auc:.4f}")

# Model comparison table
print("\n MODEL COMPARISON SUMMARY")
print("=" * 80)

comparison_df = pd.DataFrame({
    'Model': list(model_performances.keys()),
    'CV_Mean': [perf['cv_mean'] for perf in model_performances.values()],
    'CV_Std': [perf['cv_std'] for perf in model_performances.values()],
    'Test_Accuracy': [perf['accuracy'] for perf in model_performances.values()],
    'Test_Precision': [perf['precision'] for perf in model_performances.values()],
    'Test_Recall': [perf['recall'] for perf in model_performances.values()],
    'Test_F1': [perf['f1'] for perf in model_performances.values()],
    'Test_AUC': [perf['auc'] if perf['auc'] else 0 for perf in model_performances.values()]
})

comparison_df = comparison_df.round(4)
display(comparison_df)

# Find best model
best_model_name = max(model_performances.keys(),
                      key=lambda k: model_performances[k]['f1'])
best_model = model_performances[best_model_name]['model']

```

```
print(f"\n\n BEST PERFORMING MODEL: {best_model_name}")
print(f"      F1-Score: {model_performances[best_model_name]['f1']:.4f}")
print(f"      Accuracy: {model_performances[best_model_name]['accuracy']:.4f}")
print(f"      AUC: {model_performances[best_model_name]['auc']:.4f}")
```

CROSS-VALIDATION EVALUATION

Evaluating Logistic Regression...
CV Accuracy: 0.8515 (+/- 0.0714)
Test Accuracy: 0.8587
Test F1-Score: 0.8750
Test AUC: 0.8974

Evaluating Random Forest...
CV Accuracy: 0.8515 (+/- 0.0657)
Test Accuracy: 0.8913
Test F1-Score: 0.9029
Test AUC: 0.9329

Evaluating XGBoost...
CV Accuracy: 0.8502 (+/- 0.0491)
Test Accuracy: 0.8587
Test F1-Score: 0.8700
Test AUC: 0.9271

Evaluating Neural Network...
CV Accuracy: 0.8174 (+/- 0.0849)
Test Accuracy: 0.8370
Test F1-Score: 0.8515
Test AUC: 0.9134

Evaluating SVM...
CV Accuracy: 0.8556 (+/- 0.0843)
Test Accuracy: 0.8696
Test F1-Score: 0.8857
Test AUC: 0.9220

MODEL COMPARISON SUMMARY

=====

	Model	CV_Mean	CV_Std	Test_Accuracy	Test_Precision	Test_Recall	Test_F1	Test_AUC
0	Logistic Regression	0.8515	0.0357	0.8587	0.8585	0.8922	0.8750	0.8974
1	Random Forest	0.8515	0.0328	0.8913	0.8942	0.9118	0.9029	0.9329
2	XGBoost	0.8502	0.0246	0.8587	0.8878	0.8529	0.8700	0.9271
3	Neural Network	0.8174	0.0425	0.8370	0.8600	0.8431	0.8515	0.9134
4	SVM	0.8556	0.0422	0.8696	0.8611	0.9118	0.8857	0.9220

BEST PERFORMING MODEL: Random Forest
F1-Score: 0.9029
Accuracy: 0.8913
AUC: 0.9329

SECTION 8: MODEL DEVELOPMENT AND EVALUATION – ANNOTATIONS

In this section, we implement and evaluate five classification models to predict heart disease. Evaluation includes cross-validation, test performance, and selection of the best model.

1. Model Initialization

The following models are initialized with fixed random states to ensure reproducibility:

- Logistic Regression
- Random Forest
- XGBoost
- Neural Network (MLPClassifier)
- Support Vector Machine (SVM)

Each model is configured with basic hyperparameters (e.g., `max_iter=1000`, `n_estimators=100`, `eval_metric='logloss'`, etc).

2. Cross-Validation

Using `StratifiedKFold` with 5 splits, we evaluate each model using `cross_val_score` on the training set. Accuracy scores across folds are stored for comparison:

- `cv_mean` = average accuracy across folds
- `cv_std` = standard deviation of CV scores (model stability)

3. Test Evaluation

Each model is trained on the scaled training set and predictions are made on the test set:

- `y_pred` : predicted class labels
- `y_pred_proba` : predicted probabilities (if supported by model)

Metrics calculated:

- Accuracy
- Precision
- Recall
- F1-Score
- AUC (Area Under ROC Curve)

All metrics are stored for later comparison.

4. Results Summary

All performance metrics are stored in a DataFrame for easy comparison across models. The top performer is selected based on the highest **F1-Score**.

Model	CV_Mean	CV_Std	Test_Accuracy	Test_Precision	Test_Recall	Test_F1	Test_AUC
Logistic Regression	0.8515	0.0357	0.8587	0.8585	0.8922	0.8750	0.8974
Random Forest	0.8515	0.0328	0.8913	0.8942	0.9118	0.9029	0.9328
XGBoost	0.8502	0.0246	0.8587	0.8878	0.8529	0.8700	0.9271
Neural Network	0.8174	0.0425	0.8370	0.8600	0.8431	0.8515	0.9134
SVM	0.8556	0.0422	0.8696	0.8611	0.9118	0.8857	0.9220

5. Best Model: Random Forest

- **F1-Score:** 0.9029
- **Accuracy:** 0.8913
- **AUC:** 0.9328

This model offers the best balance of precision and recall, and highest discriminative power (AUC), making it the optimal choice for deployment or further tuning.

Conclusion:

The modeling pipeline evaluates various algorithms and ranks them on both cross-validation and test set metrics. Random Forest clearly outperforms others and is selected as the best baseline classifier.

SECTION 9: DETAILED MODEL ANALYSI

```
In [31]: # Confusion Matrix
y_pred_best = best_model.predict(X_test_scaled)
cm = confusion_matrix(y_test, y_pred_best)

fig, axes = plt.subplots(1, 2, figsize=(15, 6))

# Confusion Matrix
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', ax=axes[0])
axes[0].set_title(f'Confusion Matrix - {best_model_name}')
axes[0].set_xlabel('Predicted')
axes[0].set_ylabel('Actual')


# ROC Curve
if hasattr(best_model, 'predict_proba'):
    y_pred_proba_best = best_model.predict_proba(X_test_scaled)[: , 1]
    fpr, tpr, _ = roc_curve(y_test, y_pred_proba_best)
    auc_score = roc_auc_score(y_test, y_pred_proba_best)

    axes[1].plot(fpr, tpr, label=f'ROC Curve (AUC = {auc_score:.3f})', linewidth=2)
    axes[1].plot([0, 1], [0, 1], 'k--', label='Random Classifier')
```

```
axes[1].set_xlabel('False Positive Rate')
axes[1].set_ylabel('True Positive Rate')
axes[1].set_title('ROC Curve')
axes[1].legend()
axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

# Classification Report
print("\n DETAILED CLASSIFICATION REPORT")
print("-" * 40)
print(classification_report(y_test, y_pred_best,
                             target_names=['No Heart Disease', 'Heart Disease']))
```



DETAILED CLASSIFICATION REPORT				
	precision	recall	f1-score	support
No Heart Disease	0.89	0.87	0.88	82
Heart Disease	0.89	0.91	0.90	102
accuracy			0.89	184
macro avg	0.89	0.89	0.89	184
weighted avg	0.89	0.89	0.89	184

After identifying the best-performing model in Section 8, we now analyze its performance in greater detail using **confusion matrix**, **ROC curve**, and **classification metrics**.

1. Confusion Matrix

A confusion matrix is created to visualize the model's predictions on the test set:

- **True Positives (93)**: Correctly predicted heart disease.
- **True Negatives (71)**: Correctly predicted no heart disease.
- **False Positives (11)**: Incorrectly predicted heart disease.
- **False Negatives (9)**: Missed heart disease cases.

This breakdown helps identify where the model tends to make errors — e.g., more false positives than false negatives.

2. ROC Curve

The ROC curve evaluates how well the model distinguishes between the two classes:

- **X-axis:** False Positive Rate (FPR)
- **Y-axis:** True Positive Rate (TPR)
- A curve closer to the top-left corner indicates better performance.
- **AUC (Area Under Curve) = 0.933** – a strong indicator of model quality and class separation.

The dashed diagonal line represents a random guess baseline (AUC = 0.5), and our model significantly outperforms it.

3. Classification Report

This provides precision, recall, and F1-score for both classes:

Metric	No Heart Disease	Heart Disease
Precision	0.89	0.89
Recall	0.87	0.91
F1-Score	0.88	0.90
Support	82	102

- **Accuracy:** 0.89 overall
- **Macro avg:** Unweighted mean across classes
- **Weighted avg:** Mean that accounts for support of each class

These metrics confirm balanced, high-quality performance for both classes.

Conclusion:

The Random Forest model demonstrates robust classification with high AUC and strong precision/recall balance. The ROC curve and confusion matrix confirm that the model is highly effective at identifying heart disease while keeping false predictions relatively low.

SECTION 10: FEATURE IMPORTANCE ANALYSIS


```
In [32]: # Feature importance for tree-based models
if hasattr(best_model, 'feature_importances_'):
    importance_df = pd.DataFrame({
        'Feature': X.columns,
        'Importance': best_model.feature_importances_
    }).sort_values('Importance', ascending=False)

    print("Top 10 Most Important Features:")
    print(importance_df.head(10))

    # Visualization
    plt.figure(figsize=(12, 8))
    top_features = importance_df.head(10)
    plt.barh(range(len(top_features)), top_features['Importance'])
    plt.yticks(range(len(top_features)), top_features['Feature'])
    plt.xlabel('Feature Importance')
    plt.title(f'Top 10 Feature Importances - {best_model_name}')
    plt.gca().invert_yaxis()
    plt.tight_layout()
    plt.show()
```

Top 10 Most Important Features:

	Feature	Importance
11	ST_Slope_encoded	0.224569
8	ChestPainType_encoded	0.109710
3	MaxHR	0.086335
10	ExerciseAngina_encoded	0.083389
2	Cholesterol	0.083099
4	Oldpeak	0.079437
6	Exercise_Capacity	0.078521
0	Age	0.054033
1	RestingBP	0.049278
13	Chol_Category_encoded	0.035421



This section evaluates which input features had the greatest influence on the model's predictions using `.feature_importances_` (works with tree-based models like Random Forest).

- `hasattr(best_model, 'feature_importances_')`: Checks if the model supports feature importance (usually tree-based models do).
- `importance_df = pd.DataFrame(...)`: Creates a DataFrame mapping each feature to its importance score.
- `.sort_values('Importance', ascending=False)`: Sorts the features in descending order of importance.
- `print(importance_df.head(10))`: Displays the top 10 most important features numerically.

This visualization helps identify which features are driving model decisions the most — useful for model interpretability and future feature selection.

```
In [33]: # Permutation importance (works for all models)
print("\n PERMUTATION IMPORTANCE ANALYSIS")
perm_importance = permutation_importance(best_model, X_test_scaled, y_test,
                                         n_repeats=10, random_state=42)

perm_df = pd.DataFrame({
    'Feature': X.columns,
    'Importance_Mean': perm_importance.importances_mean,
    'Importance_Std': perm_importance.importances_std
}).sort_values('Importance_Mean', ascending=False)
```

```
print("Top 10 Features by Permutation Importance:")
print(perm_df.head(10))
```

PERMUTATION IMPORTANCE ANALYSIS

Top 10 Features by Permutation Importance:

	Feature	Importance_Mean	Importance_Std
11	ST_Slope_encoded	0.144565	0.021077
7	Sex_encoded	0.035326	0.006544
8	ChestPainType_encoded	0.034783	0.008134
2	Cholesterol	0.026630	0.011516
10	ExerciseAngina_encoded	0.015217	0.005325
5	CVD_Risk_Score	0.010326	0.006634
1	RestingBP	0.006522	0.004067
3	MaxHR	0.005435	0.008420
15	FastingBS	0.005435	0.006875
9	RestingECG_encoded	0.004348	0.002174

- `permutation_importance(...)` :
Computes how the model's accuracy drops when a feature's values are randomly permuted. The greater the drop, the more important the feature.
- `n_repeats=10` :
Runs the permutation multiple times to ensure stable estimates.
- `random_state=42` :
Ensures reproducibility of the results.

Result Construction

- `perm_df = pd.DataFrame(...)` :
A DataFrame with:
 - `Feature` : Feature name
 - `Importance_Mean` : Average importance score across repeats
 - `Importance_Std` : Standard deviation of importance (to judge stability)
- `.sort_values('Importance_Mean', ascending=False)` :
Sorts features by mean importance in descending order.
- `print(perm_df.head(10))` :
Displays the top 10 most influential features under this method.

Permutation importance gives a more **realistic, model-independent** understanding of feature impact — particularly valuable when using black-box models or needing to validate `.feature_importances_` results.

SECTION 11: BUSINESS IMPACT ANALYSIS

```
In [34]: # Define cost parameters
cost_false_positive = 500 # Cost of unnecessary tests/treatment
cost_false_negative = 50000 # Cost of missing a heart disease case
cost_true_positive = 5000 # Cost of early treatment
cost_true_negative = 0 # No additional cost for healthy patients

# Calculate confusion matrix values
tn, fp, fn, tp = cm.ravel()

print(f"Confusion Matrix Breakdown:")
print(f" True Negatives (Correctly identified healthy): {tn}")
print(f" False Positives (Healthy flagged as risk): {fp}")
print(f" False Negatives (Missed heart disease): {fn}")
print(f" True Positives (Correctly identified risk): {tp}")

# Calculate costs
total_cost_fp = fp * cost_false_positive
total_cost_fn = fn * cost_false_negative
total_cost_tp = tp * cost_true_positive
total_cost = total_cost_fp + total_cost_fn + total_cost_tp

print(f"\n COST ANALYSIS:")
print(f" Cost from False Positives: ${total_cost_fp:,}")
print(f" Cost from False Negatives: ${total_cost_fn:,}")
print(f" Cost from True Positives: ${total_cost_tp:,}")
print(f" Total Cost: ${total_cost:,}")
```

```
# Calculate potential savings vs no screening
baseline_cost = len(y_test) * 0.5 * cost_false_negative # Assume 50% would develop issues
savings = baseline_cost - total_cost
print(f" Baseline Cost (no screening): ${baseline_cost:,}")
print(f" Potential Savings: ${savings:,}")
print(f" ROI: {(savings/total_cost*100):.1f}%")
```

Confusion Matrix Breakdown:

- True Negatives (Correctly identified healthy): 71
- False Positives (Healthy flagged as risk): 11
- False Negatives (Missed heart disease): 9
- True Positives (Correctly identified risk): 93

COST ANALYSIS:

- Cost from False Positives: \$5,500
- Cost from False Negatives: \$450,000
- Cost from True Positives: \$465,000
- Total Cost: \$920,500
- Baseline Cost (no screening): \$4,600,000.0
- Potential Savings: \$3,679,500.0
- ROI: 399.7%

This section translates model predictions into **real-world economic costs**, helping quantify the financial value of deploying the heart disease classifier.

1. Define Cost Parameters

These are hypothetical but reasonable costs associated with each type of classification outcome:

- False Positive** = \$500
Unnecessary stress, tests, or follow-ups for someone who is healthy.
- False Negative** = \$50,000
The **most expensive mistake**: a missed heart disease diagnosis could lead to severe consequences or hospitalization.
- True Positive** = \$5,000
Cost of treating actual heart disease early.
- True Negative** = \$0
No cost when a healthy person is correctly identified.

2. Confusion Matrix Breakdown

The model's results are unwrapped into actual counts:

- True Negatives (TN)** = 71
- False Positives (FP)** = 11
- False Negatives (FN)** = 9
- True Positives (TP)** = 93

3. Cost Calculations

Each outcome is mapped to its financial cost:

- FP Cost = 11 × \$500 = \$5,500
- FN Cost = 9 × \$50,000 = \$450,000
- TP Cost = 93 × \$5,000 = \$465,000
- Total Cost of Model** = \$920,500

4. Baseline Cost (No Screening)

Assumes **50% of the total population would develop heart disease** if left unscreened:

- Baseline Cost = 184 × 0.5 × \$50,000 = \$4,600,000

5. Net Savings and ROI

- Savings = \$4,600,000 - \$920,500 = \$3,679,500
- ROI = 397.8%

This means for every 1 spent using this model, nearly 4 is saved in avoided downstream costs.

Takeaway

This analysis **clearly demonstrates the value of using the model** in a healthcare setting:

- Even though there are some false positives and false negatives, the **cost of not screening is drastically higher**.
- The high ROI shows that **machine learning-based triage systems can be cost-effective tools** in preventative healthcare.

SECTION 12: HYPERPARAMETER TUNINGS

```
In [35]: # Define parameter grids for different models
param_grids = {
    'Random Forest': {
        'n_estimators': [50, 100, 200],
        'max_depth': [10, 20, None],
        'min_samples_split': [2, 5, 10],
        'min_samples_leaf': [1, 2, 4]
    },
    'XGBoost': {
        'n_estimators': [100, 200],
        'max_depth': [3, 6, 10],
        'learning_rate': [0.01, 0.1, 0.2],
        'subsample': [0.8, 1.0]
    },
    'Logistic Regression': {
        'C': [0.1, 1.0, 10.0],
        'penalty': ['l1', 'l2'],
        'solver': ['liblinear', 'saga']
    }
}

if best_model_name in param_grids:
    print(f"🔧 Tuning hyperparameters for {best_model_name}...")

    # Create fresh model instance
    if best_model_name == 'Random Forest':
        model_for_tuning = RandomForestClassifier(random_state=42)
    elif best_model_name == 'XGBoost':
        model_for_tuning = xgb.XGBClassifier(random_state=42, eval_metric='logloss')
    elif best_model_name == 'Logistic Regression':
        model_for_tuning = LogisticRegression(random_state=42, max_iter=1000)

    # Grid search
    grid_search = GridSearchCV(
        model_for_tuning,
        param_grids[best_model_name],
        cv=5,
        scoring='f1',
        n_jobs=-1,
        verbose=1
    )

    grid_search.fit(X_train_scaled, y_train)

    print(f" Best parameters: {grid_search.best_params_}")
    print(f" Best CV F1-score: {grid_search.best_score_:.4f}")

    # Evaluate tuned model
    tuned_model = grid_search.best_estimator_
    y_pred_tuned = tuned_model.predict(X_test_scaled)

    tuned_accuracy = accuracy_score(y_test, y_pred_tuned)
    tuned_f1 = f1_score(y_test, y_pred_tuned)

    print(f" Tuned model performance:")
    print(f"   Test Accuracy: {tuned_accuracy:.4f}")
    print(f"   Test F1-Score: {tuned_f1:.4f}")
    print(f"   Improvement in F1: {tuned_f1 - model_performances[best_model_name]['f1']:.4f}")
```

```
Tuning hyperparameters for Random Forest...
Fitting 5 folds for each of 81 candidates, totalling 405 fits
Best parameters: {'max_depth': 20, 'min_samples_leaf': 4, 'min_samples_split': 2, 'n_estimators': 50}
Best CV F1-score: 0.8783
Tuned model performance:
  Test Accuracy: 0.8913
  Test F1-Score: 0.9048
  Improvement in F1: 0.0018
```

Section 12: Hyperparameter Tunings

In this section, we optimize the performance of our selected model using **Grid Search Cross-Validation** over a range of hyperparameter values.

Parameter Grids

We define specific hyperparameter grids for each model:

- **Random Forest:** number of trees (`n_estimators`), tree depth (`max_depth`), and node splitting criteria (`min_samples_split` , `min_samples_leaf`)
- **XGBoost:** includes parameters like `learning_rate` , `max_depth` , and `subsample`
- **Logistic Regression:** varies `C` (regularization), `penalty` , and `solver`

Grid Search Process

- Based on `best_model_name` , a fresh instance of the corresponding model is created.
- `GridSearchCV` is used to search across all combinations of hyperparameters using **5-fold cross-validation**, optimizing for **F1-score**.
- The best hyperparameters and their corresponding CV performance are printed.

Tuned Model Evaluation

After finding the best parameters, the model is retrained on the full training data and evaluated on the test set:

- We report test **accuracy** and **F1-score**.
- The **improvement in F1-score** compared to the previously trained model is also shown.

Results Summary

For Random Forest, the best parameters found were:

- `'max_depth': 20`
- `'min_samples_leaf': 4`
- `'min_samples_split': 2`
- `'n_estimators': 50`

Performance improved slightly:

- **Test Accuracy:** 89.13%
- **Test F1-Score:** 0.9048
- **Improvement in F1:** 0.0018

Although the gain in F1 is small, this step ensures our model is **tuned and robust** for production.

SECTION 13: MODEL INTERPRETABILITY

```
In [36]: # SHAP analysis for model interpretability
try:
    import shap

    # Create explainer based on model type
    if best_model_name in ['Random Forest', 'XGBoost']:
        explainer = shap.TreeExplainer(best_model)
        shap_values = explainer.shap_values(X_test_scaled)

        # For binary classification, take positive class SHAP values
        if len(shap_values) == 2:
            shap_values = shap_values[1]
    else:
        # Use KernelExplainer for other models
        explainer = shap.KernelExplainer(best_model.predict_proba, X_train_scaled.sample(100))
        shap_values = explainer.shap_values(X_test_scaled.sample(50))
        if len(shap_values) == 2:
            shap_values = shap_values[1]
```



```
# Summary plot
plt.figure(figsize=(20, 8))
shap.summary_plot(shap_values, X_test_scaled, plot_type="bar", show=False)
plt.title('SHAP Feature Importance Summary')
plt.tight_layout()
plt.show()

print(" SHAP analysis completed")

except ImportError:
    print(" SHAP not available. Install with: pip install shap")
except Exception as e:
    print(f" SHAP analysis failed: {str(e)}")
```

<Figure size 2000x800 with 0 Axes>



SHAP analysis completed

This section leverages **SHAP (SHapley Additive exPlanations)** to explain the predictions made by the trained model. SHAP values indicate the contribution of each feature to an individual prediction, offering both **global feature importance** and **local interpretability**.

- The code uses `TreeExplainer` for tree-based models (e.g., Random Forest, XGBoost) and falls back to `KernelExplainer` for others.
- For binary classification, only the SHAP values for the **positive class** (e.g., risk of disease) are visualized.
- The summary plot (bar type) shows the **mean absolute SHAP values** for each feature across all test samples.

Key Findings from Output:

- Only **RestingBP** and **Age** appeared in the SHAP summary bar plot. This indicates that, according to SHAP, these two features had the **greatest average impact** on model predictions.
- While this differs slightly from the feature importance rankings from tree-based models and permutation importance, SHAP provides a more nuanced perspective by considering feature interactions and individual prediction contexts.

Why This Matters:

- SHAP adds **interpretability** to otherwise opaque models (like Random Forest or XGBoost), making them more acceptable in domains like **healthcare** where decisions need to be **transparent and justifiable**.
- It helps identify which features drive model decisions, potentially uncovering unexpected influences or validating known medical knowledge.

SECTION 14: MODEL DEPLOYMENT PREPARATION

```
In [37]: # Create prediction function
def predict_heart_disease(patient_data, model=best_model, scaler=scaler,
                           feature_columns=feature_columns):
    """
    Predict heart disease risk for a single patient

    Parameters:
    patient_data: dict with patient information

    Returns:
    prediction: 0 (no risk) or 1 (risk)
    probability: probability of heart disease
    """
    # Convert to DataFrame
    patient_df = pd.DataFrame([patient_data])

    # Feature engineering (same as training)
    patient_df['CVD_Risk_Score'] = (
        (patient_df['Age'] > 55).astype(int) * 2 +
        (patient_df['Sex_encoded'] == 1).astype(int) * 1 + # Assuming M=1
        (patient_df['Cholesterol'] > 240).astype(int) * 2 +
        (patient_df['RestingBP'] > 140).astype(int) * 2 +
        (patient_df['FastingBS'] == 1).astype(int) * 1
    )

    patient_df['Exercise_Capacity'] = patient_df['MaxHR'] / patient_df['Age']

    # Select features
    patient_features = patient_df[feature_columns]

    # Scale numerical features
    patient_features_scaled = patient_features.copy()
    patient_features_scaled[numerical_features] = scaler.transform(
        patient_features[numerical_features]
    )

    # Make prediction
    prediction = model.predict(patient_features_scaled)[0]
    if hasattr(model, 'predict_proba'):
        probability = model.predict_proba(patient_features_scaled)[0][1]
    else:
        probability = None

    return prediction, probability

# Test the prediction function
print(" Testing prediction function...")

# Example patient data
example_patient = {
    'Age': 55,
    'RestingBP': 150,
    'Cholesterol': 250,
    'MaxHR': 140,
    'Oldpeak': 1.2,
    'FastingBS': 1,
    'Sex_encoded': 1, # Male
    'ChestPainType_encoded': 2, # ASY
    'RestingECG_encoded': 0, # Normal
    'ExerciseAngina_encoded': 1, # Yes
    'ST_Slope_encoded': 1, # Flat
    'BP_Category_encoded': 2, # High
    'Chol_Category_encoded': 2, # High
    'Age_Group_encoded': 2 # Senior
}

try:
    pred, prob = predict_heart_disease(example_patient)
    print(f" Example prediction: {pred} (Risk probability: {prob:.3f})")
except Exception as e:
    print(f" Prediction function needs adjustment: {str(e)}")
```

Testing prediction function...

Example prediction: 1 (Risk probability: 0.660)

Analysis and Explanation – Section 14: Model Deployment Preparation

This section focuses on preparing the heart disease prediction model for **real-world use** by creating a function `predict_heart_disease()` that accepts a single patient's data and returns both a **binary risk prediction** and the **probability** of heart disease.

What the Function Does

- **Input:** A single patient's data (as a dictionary), trained model (`best_model`), scaler, and list of feature columns.
- **Process:**
 1. Converts patient data into a DataFrame.
 2. Applies the same **feature engineering** steps used during training, including creating a `CVD_Risk_Score` and calculating `Exercise_Capacity` .
 3. Selects relevant features and **scales numerical values** using the provided scaler.
 4. Uses the model to **predict the class (0 or 1)** and **calculate the probability** (if supported).
- **Output:**
 - `prediction` : 0 (no risk) or 1 (risk)
 - `probability` : probability of having heart disease

Test Case Summary

The example input represents a 55-year-old male patient with elevated cholesterol and resting blood pressure. After passing through the prediction pipeline, the function returns:

This means the model predicts the patient is at **risk**, with a **65% probability** of heart disease.

Why This Matters

This function is the **bridge between development and deployment**:

- **Modular & reusable:** Encapsulates data prep, feature engineering, scaling, and prediction in one function.
- **Transparent:** Returns both prediction and probability, which is critical in **clinical decision support**.
- **Consistent:** Reuses the exact same transformation logic as training, reducing chances of inference-time data leakage or mismatch.
- **Healthcare-ready:** Supports a practical use case—scanning a single patient's data for diagnostic risk and returning interpretable results in real time.

Such a function is a prerequisite for production integration, such as **embedding into APIs, dashboards**, or **mobile applications** used by healthcare professionals.

SECTION 15: FINAL RECOMMENDATIONS

```
In [38]: print("\n EXECUTIVE SUMMARY & RECOMMENDATIONS")
print("=" * 60)

print(" MODEL PERFORMANCE SUMMARY:")
print(f"    Best Model: {best_model_name}")
print(f"    Accuracy: {model_performances[best_model_name]['accuracy']:.3f}")
print(f"    Precision: {model_performances[best_model_name]['precision']:.3f}")
print(f"    Recall: {model_performances[best_model_name]['recall']:.3f}")
print(f"    F1-Score: {model_performances[best_model_name]['f1']:.3f}")
if model_performances[best_model_name]['auc']:
    print(f"    AUC-ROC: {model_performances[best_model_name]['auc']:.3f}")

print(f"\n BUSINESS IMPACT:")
print(f"    Estimated Annual Savings: ${savings:,}")
print(f"    Cost per Prediction: ${total_cost/len(y_test):.2f}")
print(f"    ROI: {(savings/total_cost*100):.1f}%")

print(f"\n KEY INSIGHTS:")
if hasattr(best_model, 'feature_importances_'):
    top_3_features = importance_df.head(3)['Feature'].tolist()
    print(f"    Top Risk Factors: {' ', ' '.join(top_3_features)}")
```

```
EXECUTIVE SUMMARY & RECOMMENDATIONS
=====
MODEL PERFORMANCE SUMMARY:
  Best Model: Random Forest
  Accuracy: 0.891
  Precision: 0.894
  Recall: 0.912
  F1-Score: 0.903
  AUC-ROC: 0.933

BUSINESS IMPACT:
  Estimated Annual Savings: $3,679,500.0
  Cost per Prediction: $5002.72
  ROI: 399.7%

KEY INSIGHTS:
  Top Risk Factors: ST_Slope_encoded, ChestPainType_encoded, MaxHR
```